

Database Management Systems

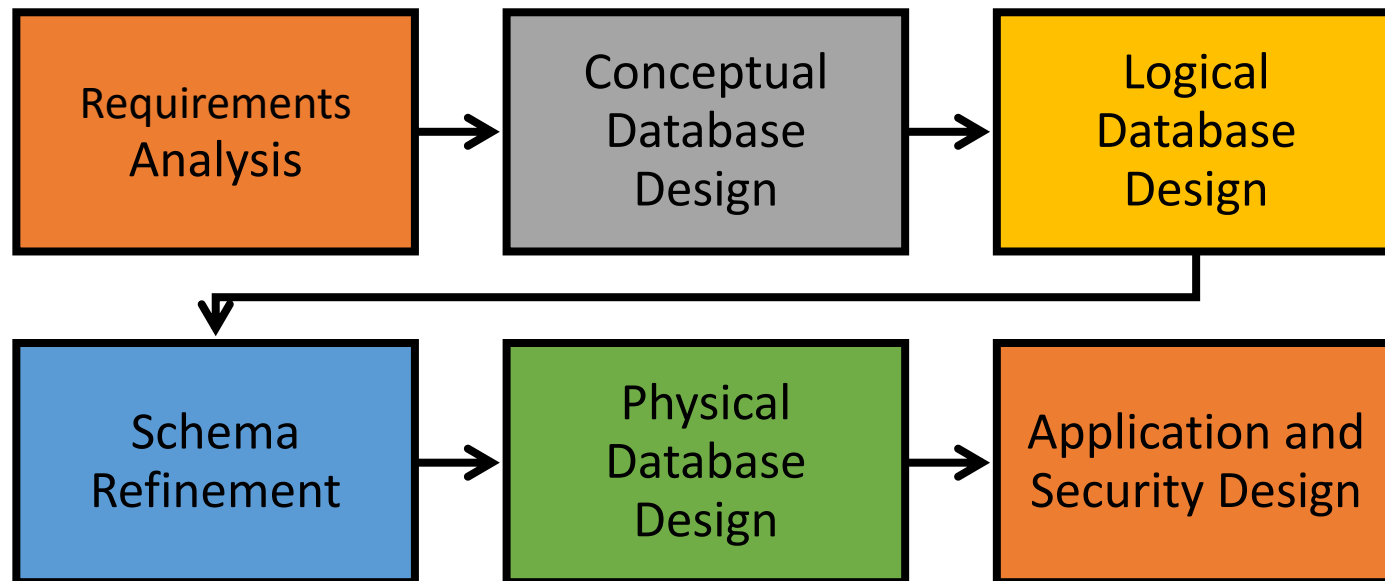
Ehsan Noei
e.noei@utoronto.ca



- Groups?
 - Send me an email (e.noiei@utoronto.ca)
 - Your names and student#
 - Until Monday (May 13) 11:59pm
 - If not,
 - Random assignments

Overview of Database Design

- The database design process can be divided into six steps:



Requirements Analysis

- What are the **entities** and **relationships** in the enterprise?
- What **information** about these entities and relationships should we store in the database?
- What are the **integrity constraints** or business rules that hold?

Conceptual Design

- The information gathered in the **requirements analysis** step is used to develop a **high-level description** of the data to be stored in the database, along with the constraints known to hold over this data.
- This step is often carried out using the **ER model**.

Logical Database Design

- Convert the conceptual database design into a database **schema** in the data model of the chosen DBMS.
 - Consider only **relational** DBMSs in this course.
 - Map an ER diagram into a relational schema.

Schema Refinement

- Normalization for desirable properties, eliminating redundancies from tables.

Physical Database Design

- Indexing, storage, etc. for performance

Application and Security Design

- We must identify the parts of the database that must be **accessible** and the parts of the database that must ***not* be accessible**.
- We must take steps to ensure that these access rules are enforced.

Entity-Relationship (ER) Data Model

- Allows us to describe the data involved in a real-world enterprise in terms of **objects** and their **relationships**.

Entity


- Real-world **object** distinguishable from other objects. An entity is described (in DB) using a set of attributes
- e.g., employee, toy department, and manager of the toy department



Employee


Attribute

- An **entity** is **described** (in DB) using a set of attributes
 - e.g., name, email, age, salary, position, etc.
- **Domain is** a set of possible values for an attribute.



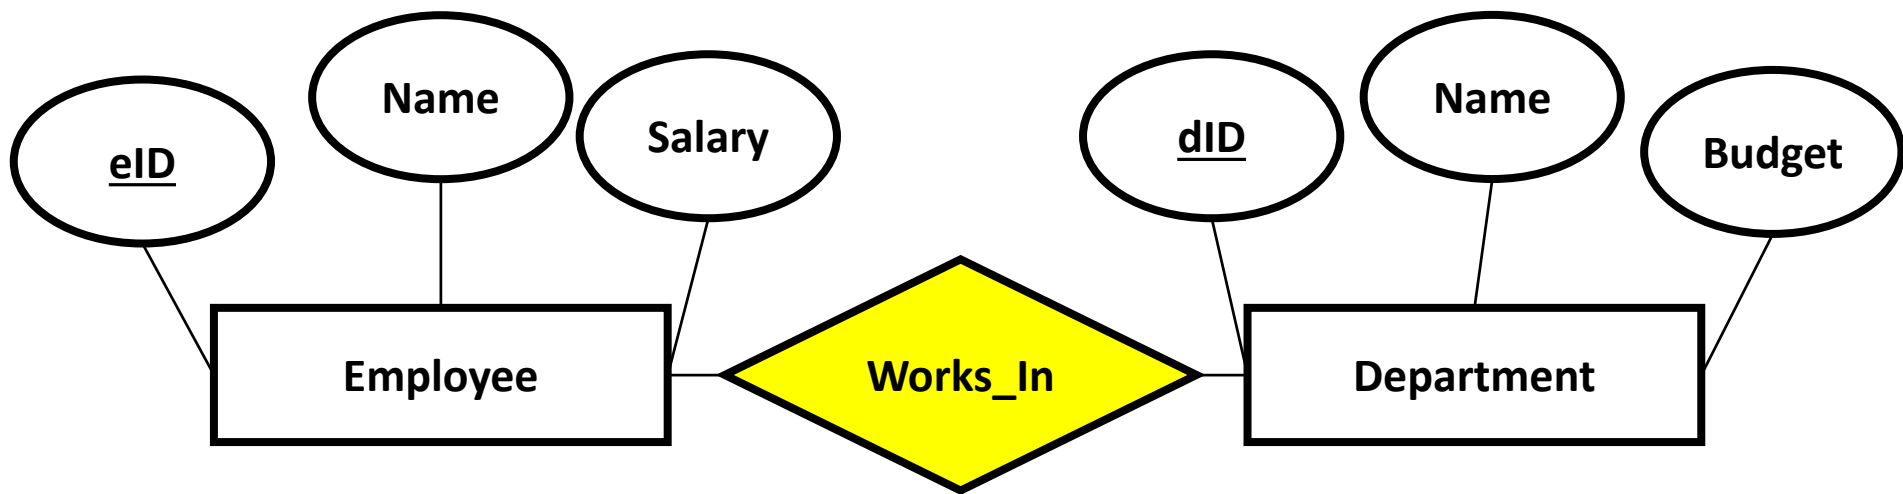
Key

- **Minimal** set of attributes whose values uniquely identify an entity
 - Many candidate key
 - Choose one of them as primary key
 - Assuming each entity set has a key




Relationship

- **Association** among two or more entities.
 - e.g., Steve works in the Toy Department




Relationship


- Relationship can also have **attributes**
 - Are used to record information about the **relationship**, rather than about any one of the participating entities




Type of Relationships




1-to-1



1-to Many



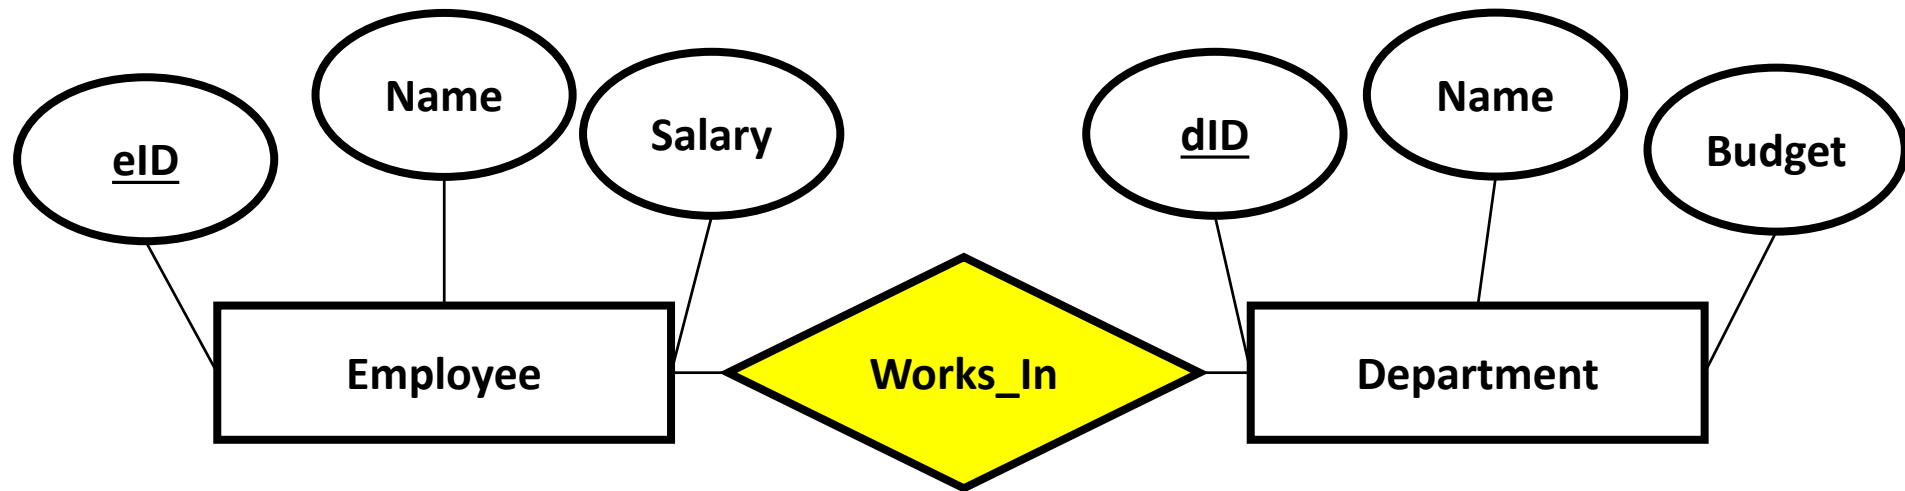
Many-to-1



Many-to-Many


Relationship Key

- All keys involved.
 - e.g., (eID,dID)




Instance of Relationship

- Snapshot of the relationship set at some instant in time.

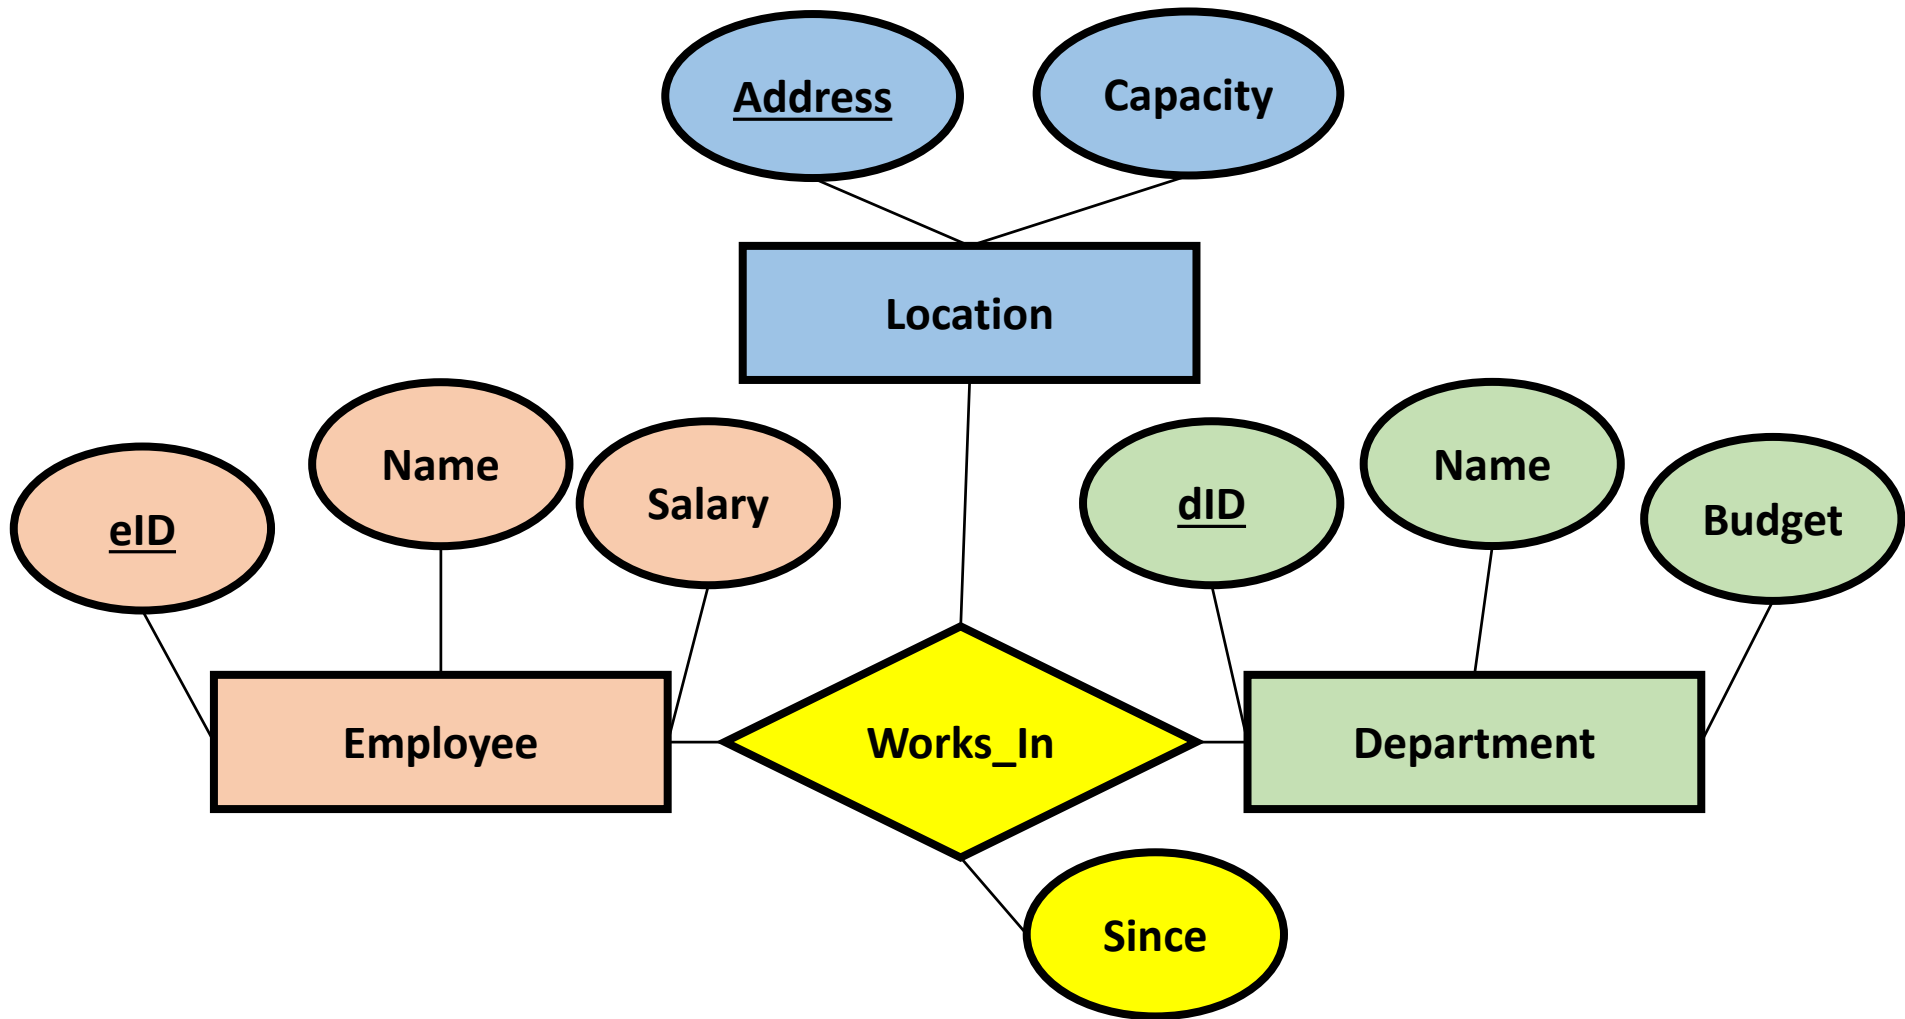


Relationship

- Suppose that each department has offices in several **locations** and we want to record the locations at which each employee works




Relationship (ternary)



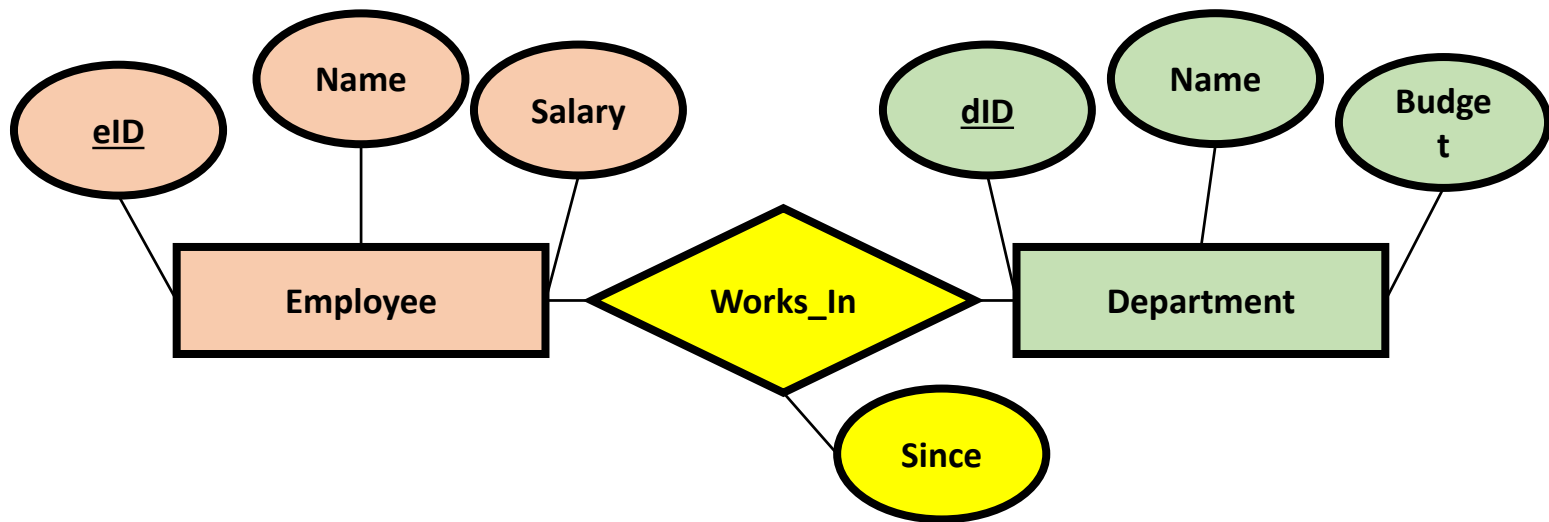
Relationship

- The entity that participate in a relationship set need **not** be **distinct**.



Key Constraints

- Consider **Works_In**
 - An employee can work in **many** departments; a dept can have **many** employees




Key Constraints

- In contrast, **each dept** has **at most one** manager, according to the *key constraint* on Manages.
- Departments are keys: given a Department entity, we can uniquely determine the Manager relationship,
 - i.e., **each department** appear **once** in the Manages table
- **One-to-many**: one employee can manage many dep., but each dep. can have only one manager

Key Constraints

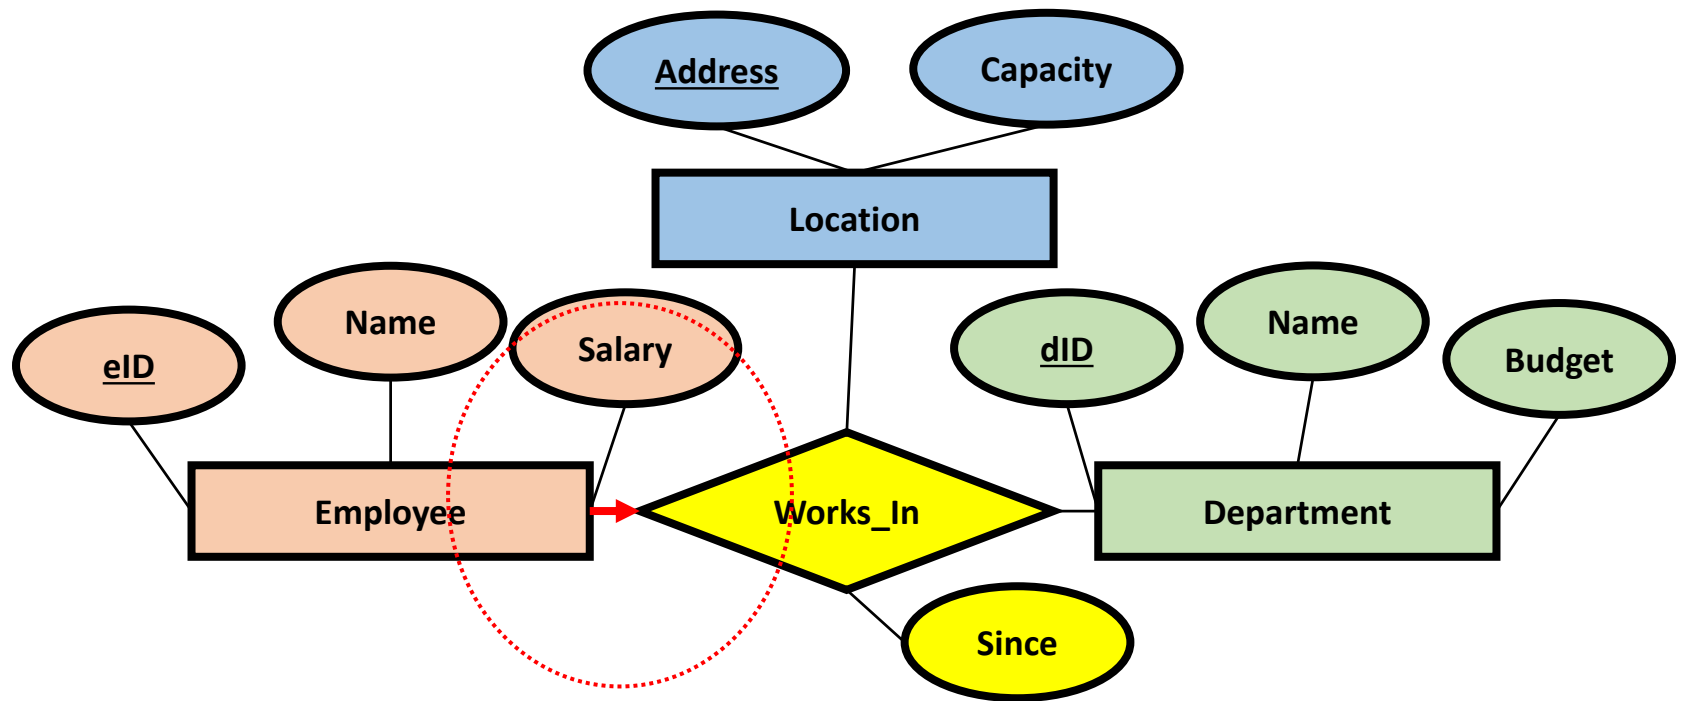
- **One-to-many**: one employee can manage many dep., but each dep. can have only one manager



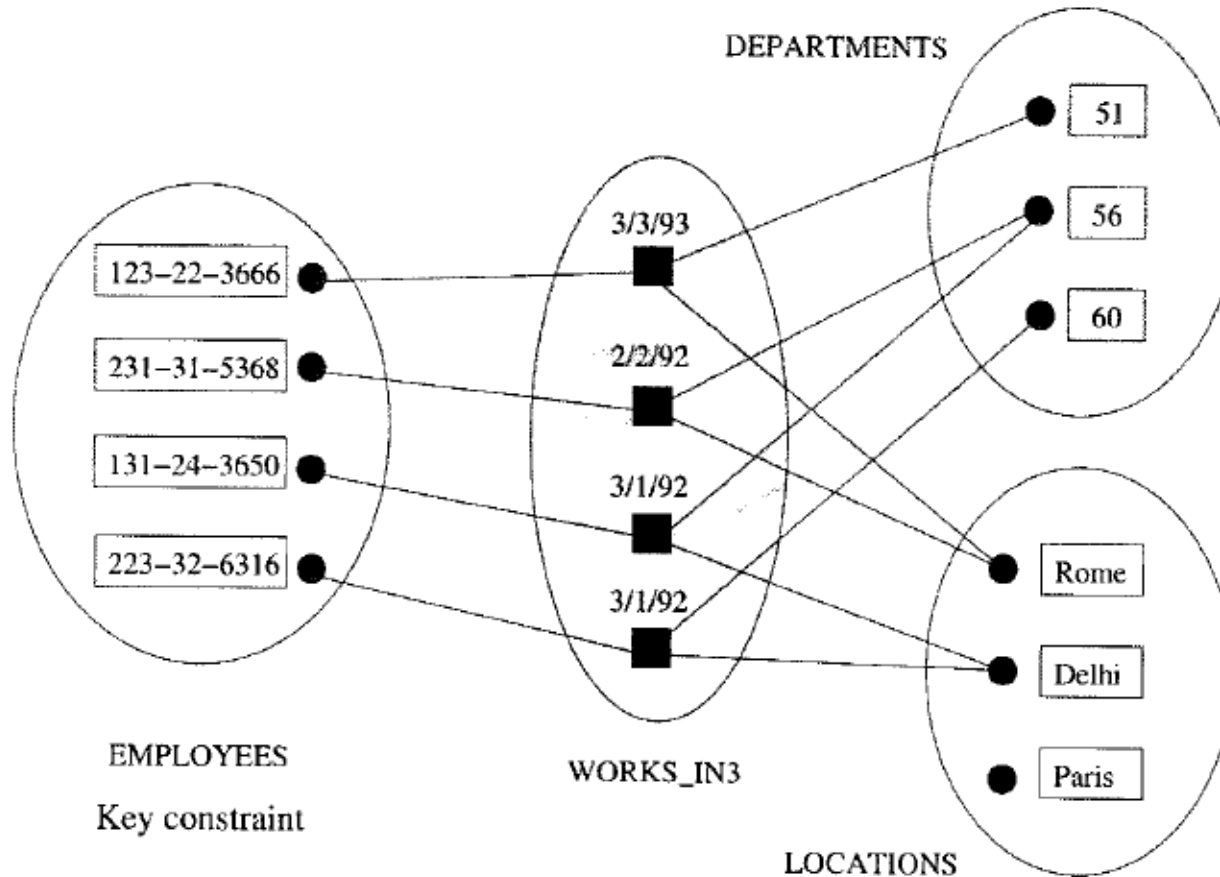
Key Constraints for Ternary Relationships

- If **each** employee works in **at most one department** and **at a single location**?

Key Constraints for Ternary Relationships



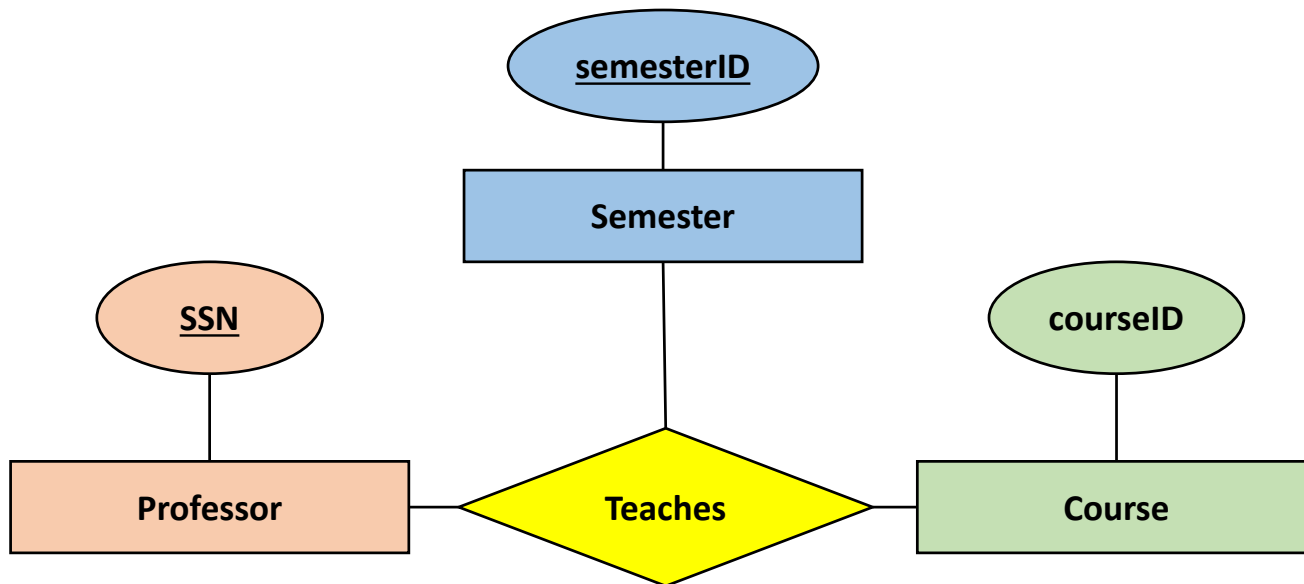
Key Constraints for Ternary Relationships



Exercise

- A university database contains information about **professors** (identified by social security number, or **SSN**) and **courses** (identified by **courseID**).
- Professors can teach the same course in several semesters, and each offering must be recorded.
- Draw an ER diagram that describes it.

Exercise




Participation Constraints

- The **key constraint** tells us that a department has **at most one** manager.
- Does **every** department have a manager?
 - If so, this is a **participation constraint**: **every department has to have such a relationship**.
 - The participation of the entity **Department** in the relationship **Manages** is said to be **total**.

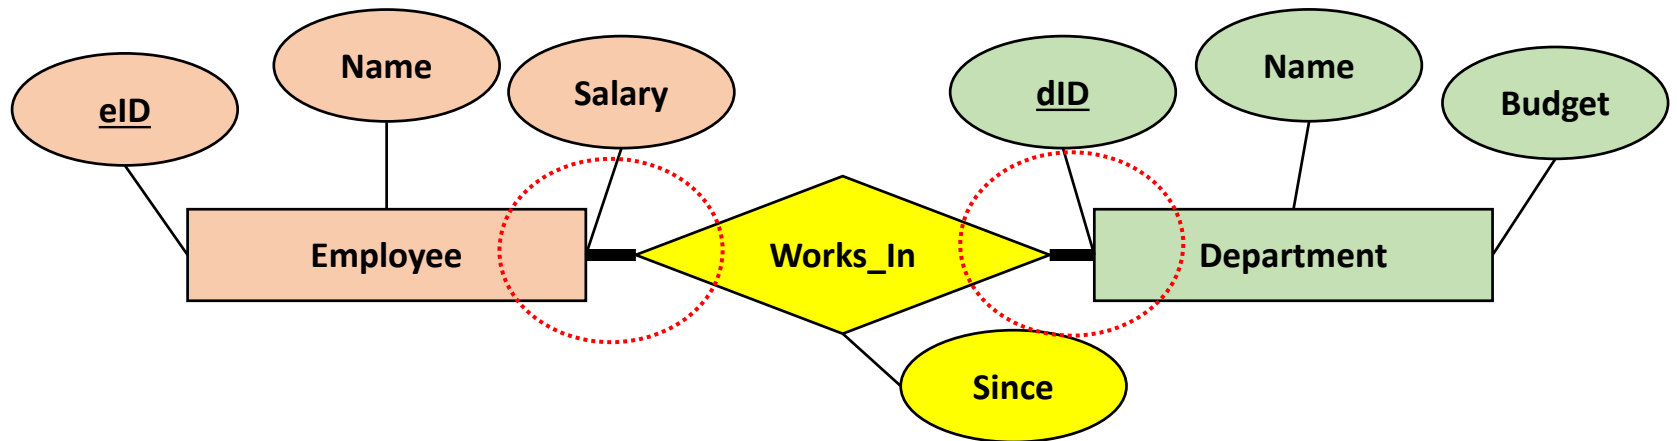
Total and Partial Participation

- The participation of the entity **Department** in the relationship **Manages** is said to be **total**.
- A participation that is **not total** is said to be **partial**.
 - not every employee gets to manage a department.

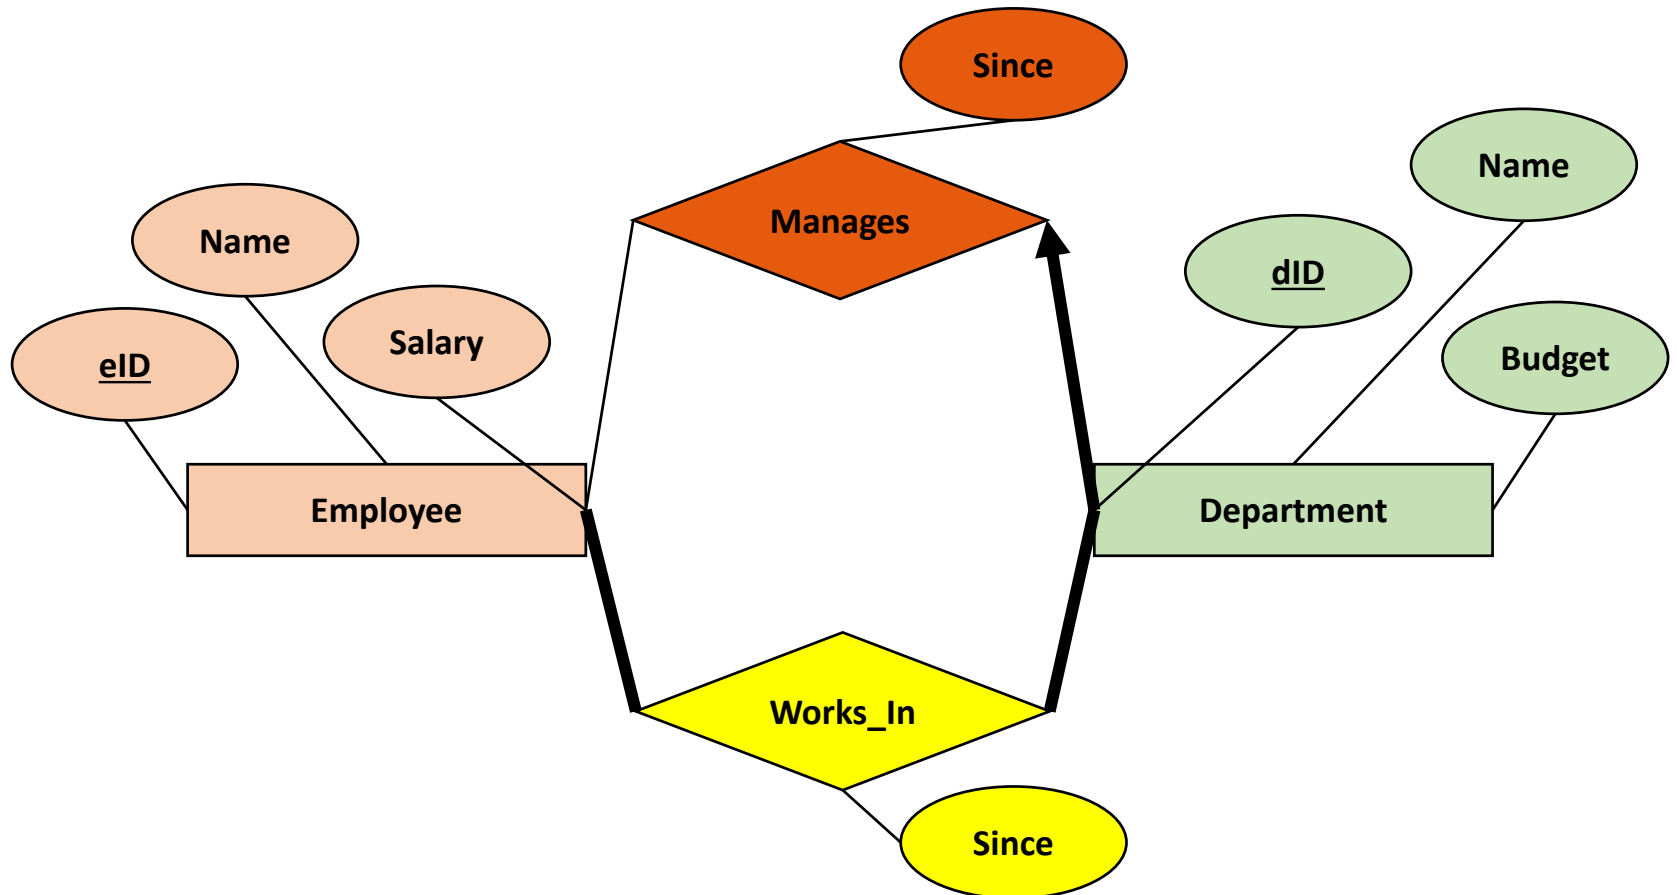


Participation Constraints

- Each **employee** works in **at least one** department and that **each department** has **at least one employee**.
- This means that the participation of both Employees and Departments in **Works_In** is **total**.



Participation/Key Constraints




Example

- A university database contains information about professors (identified by social security number, or SSN) and courses (identified by courseID).
 - 1) Every professor must teach some course.
 - 2) Every professor teaches exactly one course (no more, no less).
 - 3) Every professor teaches exactly one course (no more, no less), and every course must be taught by some professor.
- Draw an ER diagram for each that describes it.


Example

- Every professor must teach some course.




Example

- Every professor teaches exactly one course (no more, no less).



Example

- Every professor teaches exactly one course (no more, no less), and every course must be taught by some professor.



Weak Entities


- A weak entity can be identified uniquely only by considering **the primary key of another** (owner) entity
 - Owner entity and weak entity **must** participate in a one-to-many relationship set (one owner, many weak entities)
 - Weak entity set must have **total** participation
 - **Partial key**: set of attributes of a weak entity that uniquely identify a weak entity for a given owner entity

Weak Entity Example

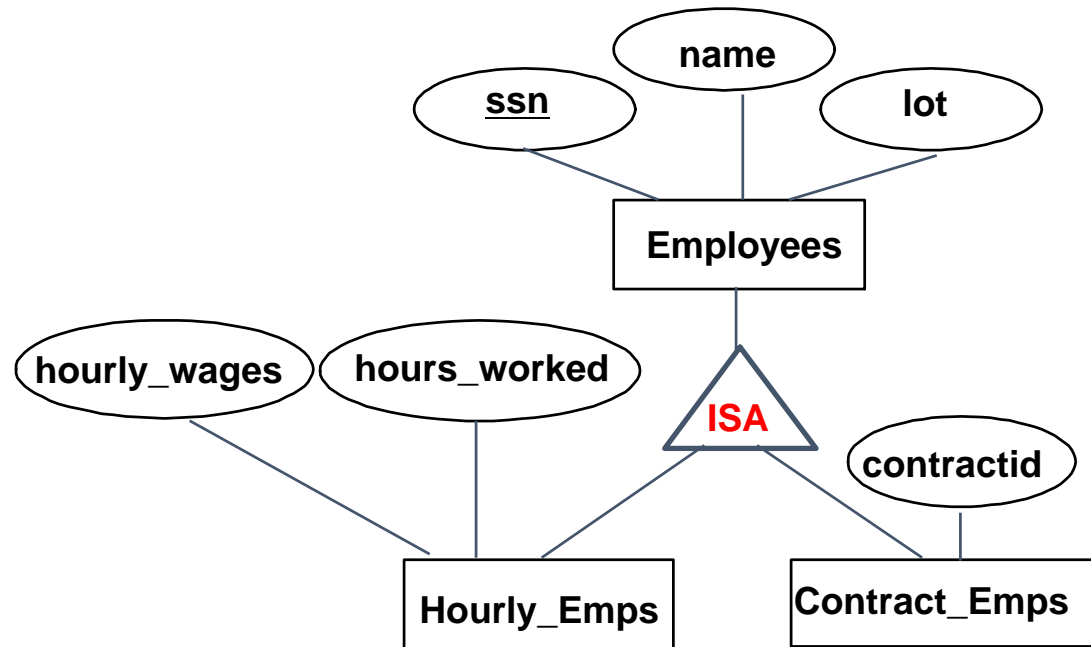
- Suppose that employees can purchase insurance policies to cover their dependents.
- We wish to record information about policies.

Weak Entity Example

- How to show? Thick borders and arrow
- Key?
 - Partial key



Class Hierarchies

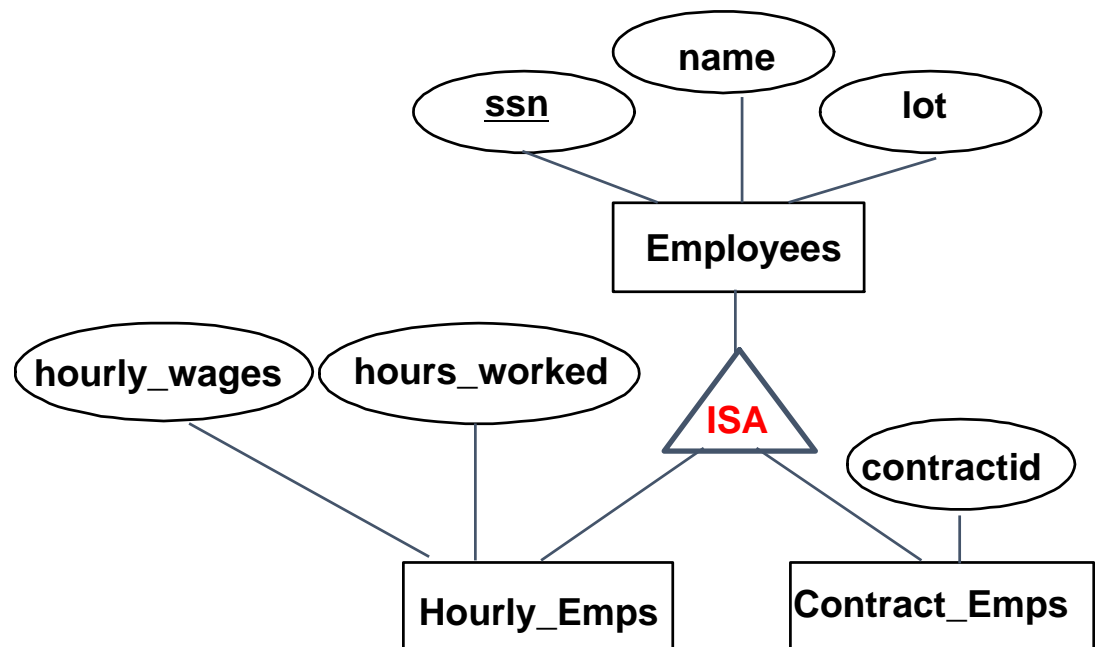


Class Hierarchies

- **Attributes** are **inherited**.
- If we declare A **ISA** B, every A entity is also considered to be a B entity.
- Constraints
 - **Overlap constraints**: Can Joe be an Hourly_Emps as well as a Contract_Emps entity?
 - No, unless stating “Contract_Emps **OVERLAPS** Hourly_Emps”
 - **Covering constraints**: Does every Employees entity also have to be an Hourly_Emps or a Contract_Emps entity?
 - No, unless stating “Hourly_Emps AND Contract_Emps **COVER** Employees”

Class Hierarchies

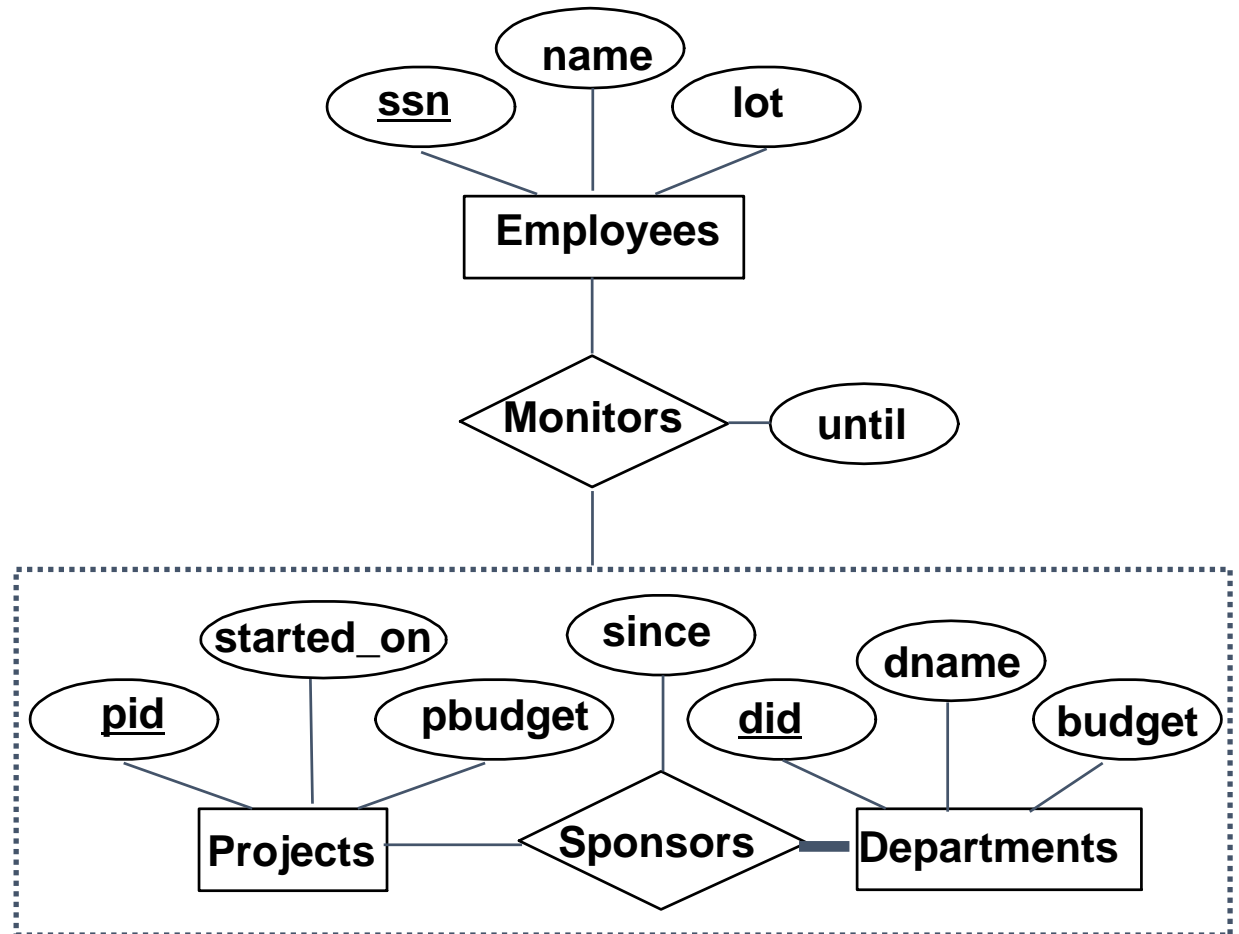
- Employees is **specialized** into subclasses.
- Hourly_Emps and Contract_Emps are **generalized** by Employees.



Aggregation

- Used when we have to model a relationship involving (entity sets and) **a relationship set**.
- Aggregation allows us to treat a relationship set as an entity set for purposes of participation in (other) relationships.

Aggregation Example

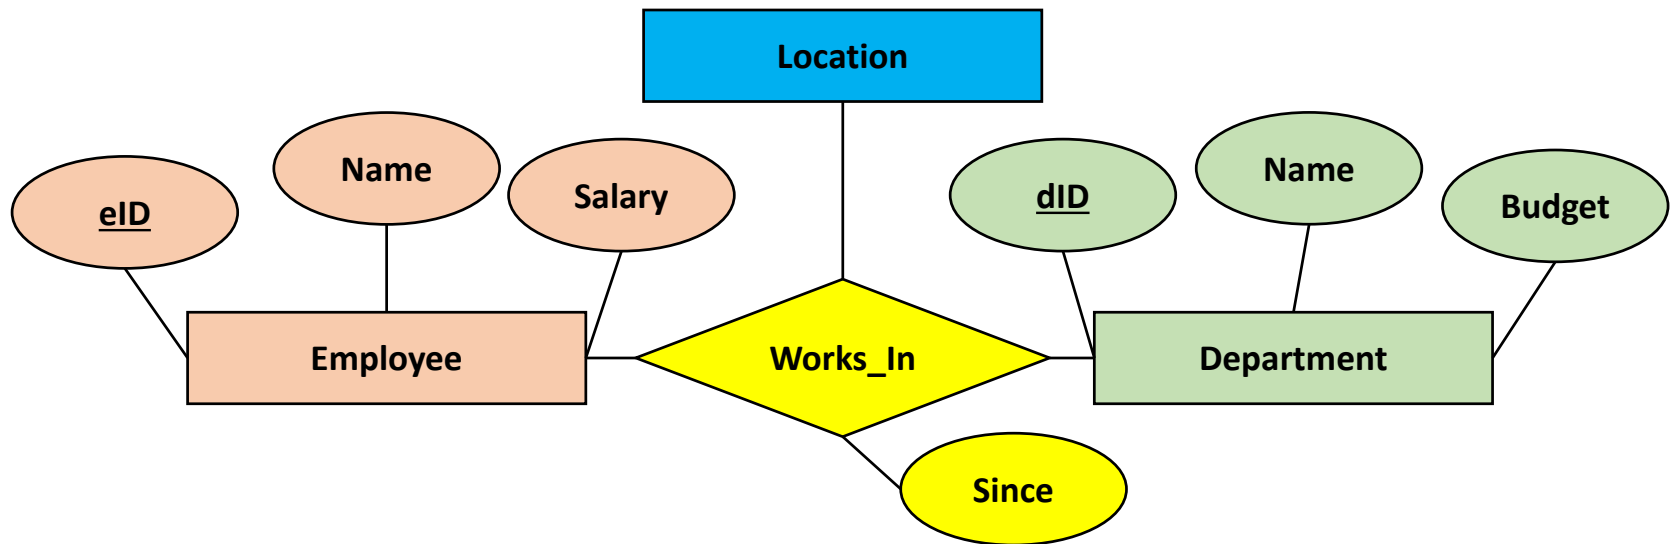


Conceptual Design Using the ER Model

- Design **choices**:
 - Should a concept be modeled as an **entity** or an **attribute**?
 - Should a concept be modeled as an **entity** or a **relationship**?
 - Identifying relationships: **binary** or **ternary**?
- Constraints in the ER Model:
 - A lot of data semantics can (and should) be captured
 - But some constraints cannot be captured in ER diagrams



Entity vs. Attribute

- We have to record **more than one** address for an employee.
- Or, we want to capture the **structure** of an address in our ER diagram.




Entity vs. Attribute Example

- Works_In **does not** allow an employee to work in a department for two or more periods
 - A relationship is uniquely identified by the participating entities
- Similar to the problem of wanting to record several addresses for an employee: We want to record several values of the descriptive attributes for each instance of this relationship.
 - Accomplished by introducing new entity set, Duration.
 - Set-valued, then attribute




Entity vs. Relationship

- Consider the relationship set called Manages.
- Suppose that each department manager is given a discretionary budget (*dbudget*).



Entity vs. Relationship

- Add *dbudget* to *Manages*?
- What if the discretionary budget is a sum that covers *all* departments managed by that employee? (redundant storage)
- Misleading; it suggests that the budget is associated with the relationship, when it is actually associated with the manager.
- Solution?
 - New entity: Managers (ISA Employee)
 - dbudget attribute of Managers



Summary of Conceptual Design

- *Conceptual design follows requirements analysis,*
 - Yields a high-level description of data to be stored
- ER model popular for conceptual design
 - Constructs are expressive, close to the way people think about their applications.
- Basic constructs: *entities, relationships, and attributes* (of entities and relationships).
- Some additional constructs: *weak entities, ISA hierarchies, and aggregation.*
- Note: There are many variations on ER model.

Summary of ER (Contd.)

- Several kinds of integrity constraints can be expressed in the ER model: *key constraints*, *participation constraints*, and *overlap/covering constraints* for ISA hierarchies.
 - Some constraints (notably, *functional dependencies*) cannot be expressed in the ER model.
 - Constraints play an important role in determining the best database design for an enterprise.

Summary of ER (Contd.)

- ER design is *subjective*. There are often many ways to model a given scenario! Analyzing alternatives can be tricky, especially for a large enterprise. Common choices include:
 - Entity vs. attribute, entity vs. relationship, binary or n-ary relationship, whether or not to use ISA hierarchies, and whether or not to use aggregation.
- Ensuring good database design: resulting relational schema should be analyzed and refined further.