

Database Management Systems

Ehsan Noei
e.noei@utoronto.ca





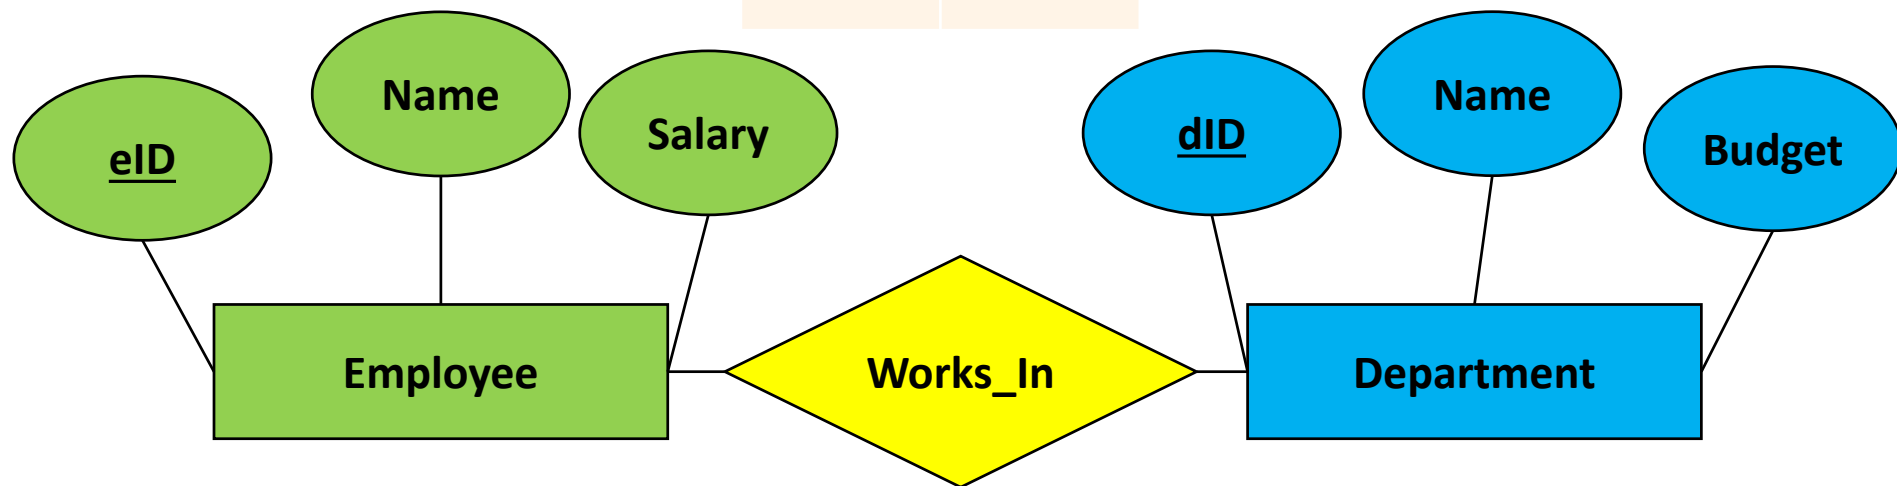
Relational Database: Definitions

- *Relation*: consists of 2 parts:
 - *Instance* : a *table*, with rows and columns. *#rows = cardinality*, *#fields = degree / arity*
 - *Schema* : specifies name of relation, plus name and type of each column.
 - E.g. Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real)

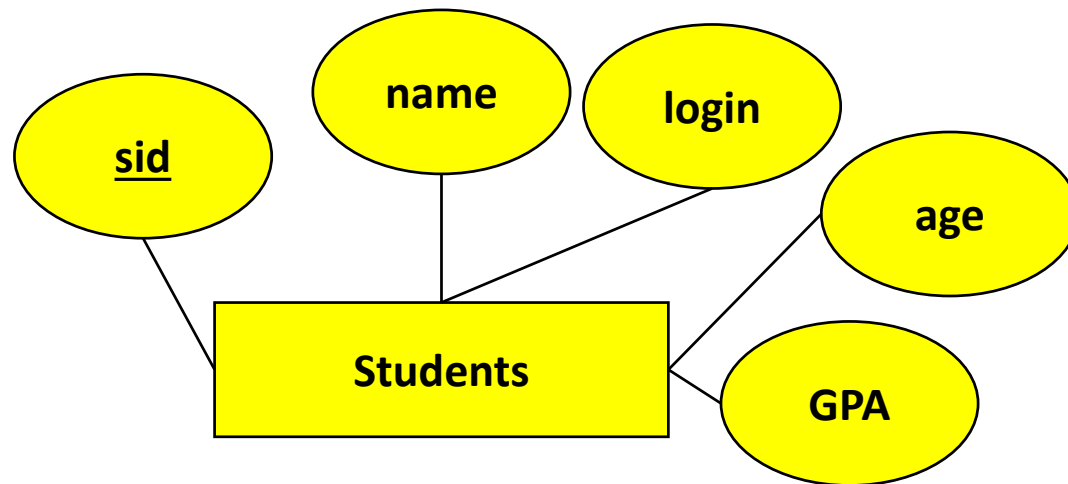
eID	Name	Salary
123	Ehsan	\$10

eID	dID
123	d20


dID	Name	Budget
d20	DB	\$20



Students (*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real)



CREATE TABLE Students (sid: string, *name*: string, *login*: string, *age*: integer, *gpa*: real)



SQL Query Language

- A major strength of the relational model: supports simple, powerful *querying* of data.
- Queries can be written intuitively, and the DBMS is responsible for efficient evaluation.
- Output of a query can be perceived as a **TABLE**.

SQL Query Language

```
SELECT < COLUMNS>  
      FROM <TABLE>  
      WHERE <CONDITION>
```


SQL Query Language

```
SELECT Name, Salary  
FROM Employee  
WHERE eID = '123'
```

eID	Name	Salary
123	Ehsan	\$10
321	Steve	\$9

SQL Query Language

```
SELECT *  
  FROM Employee  
 WHERE eID = '123'
```

eID	Name	Salary
123	Ehsan	\$10
321	Steve	\$9

SQL Query Language

SELECT *

FROM Employee

WHERE Name = 'Ehsan' AND Salary = '\$10'

eID	Name	Salary
123	Ehsan	\$10
321	Steve	\$9

SQL Query Language

SELECT *

FROM Employee

WHERE Name = 'Ehsan' OR Salary = '10\$'

eID	Name	Salary
123	Ehsan	\$10
321	Steve	\$9

SQL Query Language

```
SELECT *  
FROM Employee
```

eID	Name	Salary
123	Ehsan	\$10
321	Steve	\$9

SQL Query Language

```
SELECT E.Salary, N.Address  
      FROM Employee E, Manager N  
      WHERE E.Name = N.Name
```

Employee

eID	Name	Salary
123	Ehsan	\$10
321	Steve	\$9

Manager

mID	Name	Address
666	Joe	CA
667	Steve	CA

SQL Query Language

SELECT <COLUMNS>

FROM <TABLE>

WHERE <CONDITION>

DELETE

FROM <TABLE>


WHERE <CONDITION>

UPDATE <TABLE>

SET column1 = value1,
column2 = value2, ...


WHERE <CONDITION>

Key and Participation Constraints





```
CREATE TABLE Manages(  
  ssn CHAR(11),  
  did INTEGER,  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES  
    Employees,  
  FOREIGN KEY (did) REFERENCES  
    Departments)
```

```
CREATE TABLE Dept_Mgr(  
  did INTEGER,  
  dname CHAR(20),  
  budget REAL,  
  ssn CHAR(11),  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES  
    Employees)
```



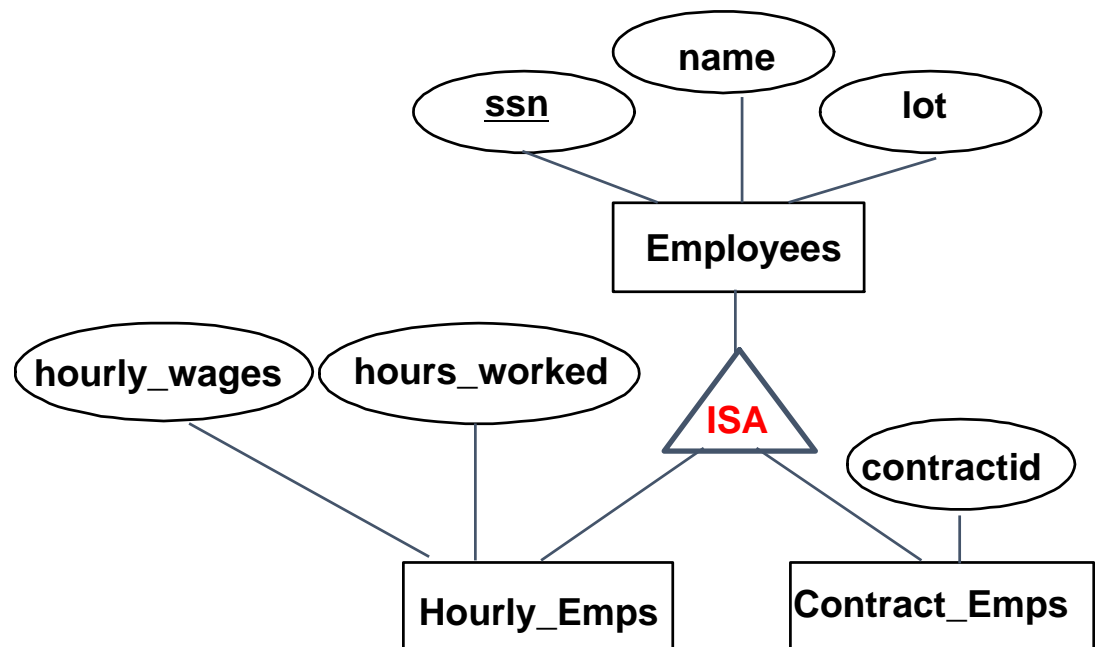
Weak Entity

```
CREATE TABLE Dep_Policy (  
  name CHAR(20),  
  age INTEGER,  
  cost REAL,  
  ssn CHAR(11) NOT NULL,  
  PRIMARY KEY (name, ssn),  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  ON DELETE CASCADE)
```



Class Hierarchies

- Employees is **specialized** into subclasses.
- Hourly_Emps and Contract_Emps are **generalized** by Employees.



Class Hierarchies

- *Overlap constraints*: Can Joe be an Hourly_Emps as well as a Contract_Emps entity?
 - *(Allowed/disallowed)*
- *Covering constraints*: Does every Employees entity also have to be an Hourly_Emps or a Contract_Emps entity?
 - *(Yes/no)*

Class Hierarchies

- **General approach:**

- 3 relations: **Employees**, **Hourly_Emps** and **Contract_Emps**.

- *Hourly_Emps*: Every employee is recorded in **Employees**. For hourly emps, extra info recorded in **Hourly_Emps** (*hourly_wages*, *hours_worked*, *ssn*); must delete **Hourly_Emps** tuple if referenced **Employees** tuple is deleted).
- Queries involving all employees easy, those involving just **Hourly_Emps** require a join to get some attributes.

- **Alternative:**

- Just **Hourly_Emps** and **Contract_Emps**.

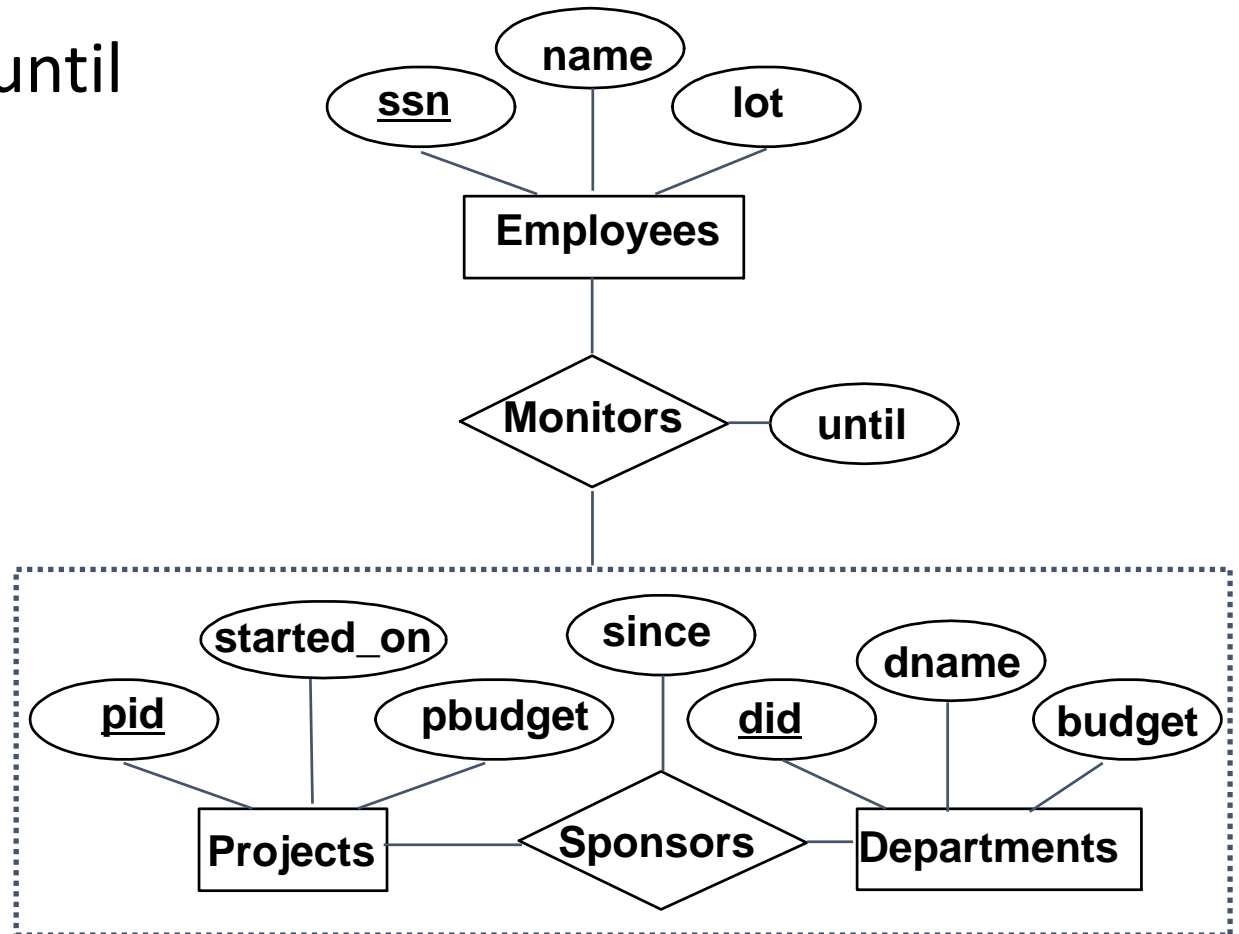
- *Hourly_Emps*: *ssn*, *name*, *lot*, *hourly_wages*, *hours_worked*.
- Each employee must be in one of these two subclasses.

Aggregation

- Used when we have to model a relationship involving (entity sets and) **a relationship set**.
- Aggregation allows us to treat a relationship set as an entity set for purposes of participation in (other) relationships.

Aggregation

- pid,did,ssn,until



Values

```
CREATE TABLE Students (  
    sid int,  
    name string,  
    PRIMARY KEY (sid)  
)
```


Values

```
CREATE TABLE Students (  
    sid int,  
    name string NOT NULL,  
    PRIMARY KEY (sid)  
)
```

Values

```
CREATE TABLE Students (  
    sid int,  
    name string DEFAULT 'sth',  
    PRIMARY KEY (sid)  
)
```

Views

- A view is just a **relation**, but **we store a definition**, rather than a set of tuples.

```
CREATE VIEW YoungActiveStudents (name, grade)
AS SELECT S.name, E.grade
FROM Students S, Enrolled E
WHERE S.sid = E.sid and S.age<21
```

Views

- Views can be **dropped** using the **DROP VIEW** command.
- How to handle DROP TABLE if there's a view on the table?
- DROP TABLE command has options to let the user specify this.

Views

- CREATE VIEW Goodstudents(sid, gpa) as select S.sid, S.gpa from Students S where S.gpa > 3.0
- How about the following:
 - INSERT into S VALUES("100", "JONE", 3.2)
 - INSERT into S VALUES("101", "Mike", 2.8)
 - DELETE from S where S.id = "100"
 - INSERT into GS VALUES("111", 3.2)
 - INSERT into GS VALUES("112", 2.8)
 - DELETE from GS where S.id = "111"

Views and Security

- Views can be used to present necessary information (or a summary), while hiding details in underlying relation(s).
 - Given YoungStudents, but not Students or Enrolled, we can find students *s* who have are enrolled, but not the *cid*'s of the courses they are enrolled in.

Summary¹

ER Model	Relational Model
Entity type	<i>Entity</i> relation
One to many or one to one relationship	Foreign key (or <i>relationship</i> relation)
Many to many relationship	<i>Relationship</i> relation and <i>two</i> foreign keys
<i>n</i> -ary relationship type	<i>Relationship</i> relation and <i>n</i> foreign keys
Simple attribute	Attribute
Composite attribute	Set of simple component attributes
Multivalued attribute	Relation and foreign key
Value set	Domain
Key attribute	Primary (or secondary) key

¹ The table is adopted from Ramez Elmasri and Shamkant B. Navathe, Fundamentals of Database Systems, Chapter 9, Sixth Edition